STRATA: A Process-Mining Framework for Detecting and Locating HVAC Faults from Building Trend Data

Yassh Singh
Thompson Rivers University, Kamloops, Canada
Undergraduate Research Experience Award Program (U-REAP) 2026
Supervisors: Dr. Anthony Aighobahi and Dr. Mridula Sharma
singhy21@mytru.ca

Abstract—Faults and poor control in heating, ventilation and air-conditioning (HVAC) equipment waste 15 to 30 percent of a commercial building's energy, and many go unnoticed. Building automation systems log the operating data and display it as trend graphs, so a fault is found only when a technician reads the graphs. Automated tools either apply rules written in advance, which miss any fault that was not anticipated, or report anomaly scores that say a day is unusual without saying where or what is wrong. This paper presents STRATA, a framework that works from the trend data a building already records and learns, from one period of normal, fault-free operation, what each day should look like. It then examines every new day with eight checks, each asking whether a physical relationship or an operating sequence has moved outside its normal range. Each check is first run on fault-free days it has never seen. The number of those days it wrongly flags is its false-alarm rate. A check may then raise an alarm only when it flags far more days than that false-alarm rate could explain. Every alarm states what broke, where, the evidence and the check's measured false-alarm rate. Process mining supplies the structure. The sensor data becomes a log of events along the control sequence, and models of normal operation are discovered from that log for each zone's equipment and for the whole unit, so every alarm knows which piece of equipment it belongs to. On five public benchmark systems from Lawrence Berkeley National Laboratory, each simulated with faults whose cause and location are known, STRATA detected 158 of 175 faults, most on the first day, and raised false alarms on 26 of 480 fault-free days it had never seen. In the 50 cases where the fault lay inside one zone's equipment, the alarm pointed to the correct zone in 37, to no zone in 13, and to a wrong zone in none. Moving the framework to a new building means writing one configuration file, not new code. The planned next step is a live trial on a real building. It is also the first step toward ProcessHeal, a self-healing building that finds its own faults, locates them, works out their cause and corrects them. STRATA does the finding and the locating; working out the cause and correcting it are the later stages.

Index Terms—HVAC, fault detection, process mining, building automation, event logs, explainable alarms

I. Introduction

Buildings account for about a third of energy-related carbon emissions [1], and HVAC accounts for up to half of a building's energy use [2]. Faults and poor control that go uncorrected waste 15 to 30 percent of a commercial building's energy [3]. A study of more than 60,000 pieces of HVAC equipment found 40 percent of air handlers in a faulted state on any given day, in buildings that already had automation systems with alarms [4]. Seven of the ten most common air-handler faults in that study are faults of behaviour, such as a wrong schedule or an unreleased manual override, rather than a broken part.

A building operator today usually learns of such a fault in one of three ways. An occupant complains that a room is too hot or too cold. The operator notices that a line on a trend graph has drifted from where it usually sits. Or an automated detector raises an alarm. The first two find the fault late; the third works only as well as the method behind it.

Those automated methods fall into two groups. Rule libraries, which come built into commercial building automation systems, are transparent and need no training data, but they catch only the faults for which someone wrote a rule. Machine-learning detectors learn from the building's own data and so can find faults that no rule anticipated, but most report an anomaly score: a number that says today is unusual without saying where or what is wrong. A score is not enough to write a work order.

This paper presents STRATA. It is configured with rules, written once per building in a single file, but those rules do not raise alarms. They turn the sensor data into a log of events: cooling started, damper open while commanded closed, a supply-air reading that physics says cannot be right. What a normal day of those events looks like is then learned from a period of the building's own fault-free operation, and every new day is checked against it, so a fault can show up as a departure from normal even when no rule names it. Each alarm states the fault, the piece of equipment, the evidence, and how often that same check has been wrong on fault-free days.

The method behind it is process mining, which discovers a model of a process from a log of events and checks later behaviour against the model [5]. STRATA builds such a log from the sensor data, for each zone's equipment and for the whole unit, which is how an alarm knows where it belongs. To our knowledge this is the first use of process mining for HVAC fault detection (Section II).

On five public benchmark systems with 175 seeded faults, STRATA detected 158, most on the first day, and placed no alarm in a wrong zone (Section IV). The code and the complete evaluation are public.[1]

II. Related Work

A. Rule-based detection, the method in use today

Fault detection in real buildings is almost entirely rule-based. The air-handling-unit performance assessment rules of Schein and colleagues set the pattern [6]. A rule states a relationship that must hold in a given operating mode, such as no water flow through a closed valve, and fires when it does not. ASHRAE Guideline 36 publishes such fault conditions alongside its control sequences [7], [8], and an open-source implementation of them exists [?]. A survey of fourteen commercial fault-detection tools found this rule-and-threshold style to be the norm [9], and interviews with 29 industry experts in 2024 found it still so [10].

Detection pays when it is used. Among 26 organisations running fault-detection software on 550 buildings, the median energy saving was 8 percent [11]. In a two-site field trial of Guideline 36, the facility team paid no attention to the built-in fault detection [12]. A 2026 California code study names alarm overload and a shortage of trained staff as the main barriers [13], and when alarms outnumber the people who can act on them, false alarms erode trust in the true ones [14].

Rules work from the first day and need no baseline data. Their limit is structural: they miss any fault that no rule author anticipated. The Lawrence Berkeley National Laboratory (LBNL) group itself notes that commercial diagnostics often list several candidate causes or none [15], and recent work answers that with active tests that command the equipment and watch the response [16]. The same group has carried the rule approach through to correction, field-testing seven hand-coded correction algorithms on 225 pieces of equipment in four buildings [17], [18].

B. Machine learning, and what its alarms leave out

A 2024 systematic review covers 211 studies of machine learning for HVAC fault detection [2], and a 2025 review follows the field through the equipment lifecycle [19]. On the standard benchmarks, accuracy has saturated at 97 to 98 percent [20], [21]. What has not been solved is the output. Where an explanation is offered, it comes from a separate layer that ranks the sensors by how much each one swayed the verdict [22], [23], [24]. That answers which signal mattered, not what the equipment should have been doing. Two 2026 studies come closest to a located alarm. One estimates the position of a stuck damper or valve after training on fault-free data alone [25]. The other weighs the evidence for each component that could be at fault and names the most likely one, without labels [26].

The second problem is labels. Real buildings have none, so the strongest recent work trains only on normal operation and flags whatever departs from it [27], [28], and what it reports is a single deviation score. A model learned on one building also travels badly to another: in a 2026 study, detectors lost almost half their performance on buildings they had not seen [29], [30], [31]. A live deployment on a Danish school building chose its detector for the fewest false alarms [32], and an audit of 65 such studies found two that shared their code [33].

Two recent hybrids use hand-written rules to label historical data and then train a classifier on those labels: one on a Danish university dataset [34], the other on a real air handler in Norway with the Schein rules as its rule layer [35]. The rules are still written in advance, so the classifier can learn only the faults the rules already name, and the output is still a class label without a location. A newer direction has a large language model write and refine the rules under physical constraints and report its evidence; it beats the neural detectors on the LBNL air-handler data [36], although a 2026 review of 66 such studies found none in sustained deployment [37].

C. Process mining, and how close it has come to buildings

The discovery algorithm used here, the inductive miner, is the standard one [38]. The check lines a day's events up against the model and shows where they depart from it [39], so the explanation comes with the method rather than from a layer added afterwards. Applying it to sensor data [40] needs one extra step, event abstraction, which turns continuous signals into discrete events [41], [42]. How to do that well, and how to keep the resulting events checkable by a person, is an open question in the process-mining community as of 2026 [43], [44].

With that step the method has reached several physical systems. It detects anomalies in chemical plants [45], in industrial alarm sequences [46] and in industrial sensor streams [47]. One group has carried it to digital twins and to fault diagnosis on a robotic arm, and shows that features drawn from the process model make a detector far better than the plain check alone [48], [49], [50]. The closest analogue to a building is pasteurisation, where sensor signals are abstracted into activities and checked against the critical control points of the HACCP food-safety standard [51]; the parallel between those control points and Guideline 36 control sequences is the one STRATA builds on. Two kinds of work sound similar but are not this. Petri nets of HVAC equipment drawn by hand exist [52], but a net an engineer draws is not a model discovered from a log. And LBNL's OpenBuildingControl verifies that a control implementation conforms to its Guideline 36 specification before the building runs [53]; STRATA checks what the building does afterwards, from its operating data.

[1] https://github.com/iyassh/strata

D. The gap, and where STRATA sits

No published work applies process mining to building HVAC fault detection. The 211-study review has no process-mining entry in its taxonomy of 23 methods [2], and the LBNL taxonomy of fault-detection methods does not list it [9]. From the other side, a review of 47 process-mining studies in Industry 4.0 [54], a review of 36 on sensor data covering industry, smart homes and healthcare [42], and a review of 93 on process management with the Internet of Things [55] contain no building application. Eight Google Scholar searches in February 2026 found no such paper, and a title-and-abstract search of the OpenAlex index for process mining together with HVAC, air handling, building automation, building energy or chillers, run on 16 October 2026, returns no work since 2023. The programmes of the process-mining conferences and workshops of 2024 to 2026 contain none. The nearest work analyses building-management alarm logs statistically, without discovering a process model from them [56].

The gap is not hard to explain: the two communities do not meet, and buildings log continuous sensor streams rather than discrete events, so the event-abstraction step of Section II-C stands in the way.

Katipamula and Brambley divide the problem into detection, isolation and identification: noticing that a fault exists, locating it, and naming its cause [3]. Only the two 2026 studies above reach isolation. STRATA does detection and isolation, and sits between the two groups: it learns normal operation from fault-free data, which a rule library cannot, and reports a location, which a score cannot. It is a complement to a rule library rather than a replacement, and its located alarms are the kind of input the LBNL correction algorithms were built to act on. Identification of the cause, and correction, are the later stages of ProcessHeal, the self-healing building toward which STRATA is the first step.

III. The Framework, Stage by Stage

Figure 1 shows the framework as four tiers. The top two tiers are built once per building from one fault-free period; the bottom two run on every new day. The stages below follow the figure from top to bottom.

A. Stage 1: the building's data and one configuration file

The input is the trend data a building-automation system already records, usually once a minute: temperatures, flows, pressures, damper and valve positions, fan speeds. No new sensors are needed. The framework needs one fault-free period of that data, ideally a year so that every season is seen, and one configuration file. That file is the only thing specific to a building. It does two things. First, it maps the building's sensor names to the roles the framework understands, so that a point known in the building only by a code is known to the framework as, say, the supply-air temperature. Second, it writes down how the equipment is supposed to behave, as rules. The rules come from ASHRAE Guideline 36 [7], the industry's published standard for how HVAC equipment should be controlled, and from the equipment's own documents, which often state what its readings should be. A few more rules we wrote ourselves, later, after the models of Stage 3 showed us things that fault-free days always do. There are two kinds of rule. State rules describe what the unit is doing: whether it is occupied, whether it is cooling, whether it is bringing in outside air. Signature rules describe what must never happen, such as flow through a closed valve. In Stage 2 these two kinds of rule produce the two kinds of event shown in Figure 1. Onboarding a new building means writing this file, not changing code. On the benchmarks, configuring a second equipment type took about one and a half hours, and two of the five main systems were configured before any fault data was opened, so nothing in them could have been tuned to the faults.

B. Stage 2: turning sensor data into events

Process mining does not read sensor values. It reads a log of events, so the first thing the framework does with each day of data is turn it into one. The rules of Stage 1 do this. Every minute, the state rules are applied to the readings, and whenever the state of the equipment changes, an event is written with its time: "fan started at 06:02", "cooling became active at 09:14". These are state events. They say what the equipment is doing and pass no judgement on whether it is right. The signature rules are applied in the same way, and whenever one is broken an event is written: "water flowing through a closed valve at 09:40". These are signature events. They say that something happened which should not. Each day becomes one sequence of events. The two kinds of event are kept apart from here on. They do two different jobs. A signature event catches a fault that someone named in advance. The model catches a day that departs from normal in a way nobody named. State events go to Stage 3, where the model of a normal day is discovered from them. Signature events never do; they go straight to the Rules check of Stage 4, the dashed arrow in Fig. 1. The reason is circularity. If the model were discovered from a log that contained signature events, it would learn that fault-free days have none, and on a fault day it would flag the signature event as a departure from normal. That would look like process mining finding a fault, but it would only be our own rule being counted a second time. We measured this on one system: letting signature events into the model raised the unit conformance check's apparent detection from 1.8 to 5.2 percent of fault days, and every one of the added days was the model re-finding a rule we had written ourselves.

C. Stage 3: discovering what normal looks like

This is the process-mining step. Its input is the state events of every fault-free day, from Stage 2. Its output

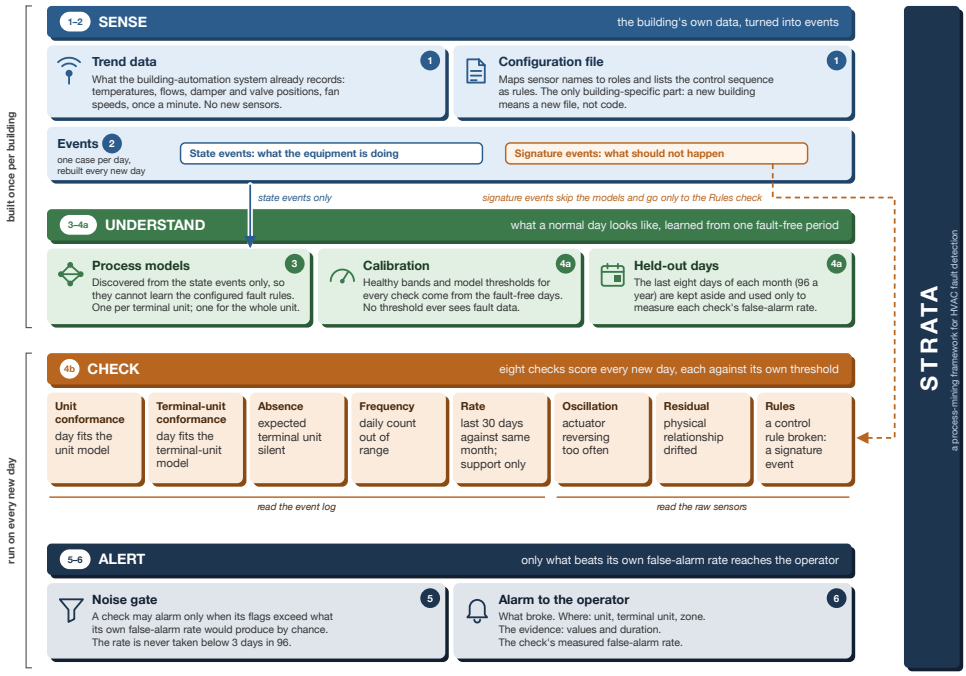


Fig. 1. The STRATA framework as four tiers. The top two are built once per building from one fault-free period; the bottom two run on every new day. The badges number the stages; 4a and 4b are the two halves of Stage 4. The arrows carry the rule that matters most: state events feed model discovery; signature events never do, and go only to the Rules check (Stage 2 says why).

is a model of a normal day. A model is a small map of the states the equipment passes through and the order it passes through them in. For an air handler it might read: the fan starts, the occupied period begins, cooling becomes active, and later in the day each of these ends. Nobody draws this map. A standard process-mining algorithm, the inductive miner [38], reads all the fault-free days and finds the one map that they all follow. Process mining calls this map a Petri net. The framework discovers models at two levels, and these are the two strata that give STRATA its name. The first level is the unit, meaning the main piece of equipment, such as an air handler. Its model is the normal daily cycle of that equipment. The second level is the terminal unit, meaning the box in each zone that the air handler serves. Each box has a damper and a reheat valve, and each box gets its own model of the normal day of those two parts. This is what gives an alarm its location. In Stage 4 a new day is compared against each model in turn. If the day fits the air handler's model and fits every terminal unit's model but one, the alarm names that one terminal unit. Reading the models during development also showed patterns of normal operation we had not thought to write down, such as the heating never running on an occupied workday; those became the extra rules of Stage 1.

D. Stage 4: calibration, then eight checks

Calibration (4a). Before any new day is scored, the framework learns from the fault-free period what normal looks like for this building, as numbers. Most checks work the same way. A check looks at one day and reduces it to one number, such as how many times the heating started. On fault-free days that number stays within some range, and that range is the check's healthy band; its edges are the check's thresholds. A day whose number falls outside the band is flagged. The last eight days of each month, 96 days in a year, are held out and used for one thing: counting how often each check flags a fault-free day, which is its false-alarm rate. The other days set the bands and train the models. The two conformance checks set their band on a further set of fault-free days, kept apart from both, as Section VI explains. No band ever sees fault data, and none sees the days on which its false-alarm rate is later reported.

The checks (4b). Every new day is scored, meaning that all eight checks are run on it, and an alarm names the check that raised it. The first three read the raw sensor values. The Rules check flags a day on which a signature event of Stage 2 occurred, such as flow through a closed valve. It is the one check that needs no band, because the rule itself says what is wrong. The Residual check uses a relationship that physics says must hold in a known state. With the cooling coil off, for example, the supply-air temperature should equal the mixed-air temperature plus the small rise the fan adds. The gap between the two sides is the day's number, and the day is flagged when that gap leaves the band. The Oscillation check counts how many times an actuator, such as a damper,

```
STRATA alarm -- series fan-powered terminal
               units, 2018-01-01 (day 1 of fault)
 what      zone primary-air flow not tracking
          its setpoint          [check: rules]
 where     terminal unit TU_S, zone S
 evidence  715 CFM measured against a 200 CFM
          setpoint, a 515 CFM excess, held for
          720 min while occupied; damper 0.80
          open, 0.32 on the same date of the
          fault-free year
 rate      rules check: 0 false-alarm days on
          the 96 held-out fault-free days
```

Fig. 2. One alarm from the series fan-powered benchmark, every value taken from the framework's recorded output. The seeded fault was the zone-S damper stuck at 80 percent: location and symptom correct, caught on day 1.

reversed direction during the day. A count far above the fault-free band is the sign of an unstable control loop. The other five read the event log of Stage 2. The Unit conformance and Terminal-unit conformance checks take the day's state events and measure how well they fit the models of Stage 3, one against the unit's model and one against each terminal unit's model. A day fits poorly when, for example, cooling starts before the fan, an order the fault-free days never showed. The Absence check looks for a terminal unit that produces events on every fault-free occupied day and produced none on this occupied day. The Frequency check counts the events of each kind in the day, such as the number of heating episodes, and flags a count outside the band, whether too many or too few. The Rate check looks over a longer window. It counts on how many of the last 30 days a state, such as heating, was active, and compares that with the same month of the fault-free year. This check never raises an alarm on its own, because on one system it added false alarms. It is reported beside another check's alarm as supporting evidence.

E. Stage 5: the noise gate

A check may raise an alarm on a scenario only when the number of days it flagged is too large to be explained by its own false-alarm rate. The test is a binomial test, the calculation used for coin tosses. If a check wrongly flags three of 96 held-out days, flagging forty of 365 is far beyond chance and the check may alarm; flagging five is not. The rate is never taken below three in 96, so a check with no false alarms is not treated as perfect; the alarm itself reports the measured count, which can be zero. On the benchmarks the test is applied once per check per scenario, and the first day the credited check flagged is the day of the alarm.

F. Stage 6: the alarm the operator receives

What survives the gate reaches the operator with four things: what broke, where, the evidence, and the check's measured false-alarm rate. Figure 2 shows one.

IV. Results on Public Benchmark Data

We evaluated the framework on the public HVAC fault-detection datasets from LBNL [57]. Each simulated dataset is one-minute data for one equipment type, recorded once fault-free and once per seeded fault: a stuck damper, a biased sensor, a fouled coil, a leaking valve, an unstable controller. Table I lists the seven datasets, one of them a measured year from a real rooftop unit. We found no published detection results on the fan-powered or dual-duct datasets. Unless stated, the results are for the five main systems, which have a full year of data. The simulated rooftop unit alone records the refrigerant side; the measured unit has 182 fault-free days and one fault case, so it yields a false-alarm rate but no detection count.

TABLE I
The seven LBNL datasets. "Scored" is the number of seeded faults evaluated; "Detected" those the framework caught; "False-alarm days" the held-out fault-free days on which it alarmed. †One of these was attributed on audit to a dataset defect rather than to the fault; the totals in the text count it, and are 157 without it.

System	Scored	Detected	False-alarm days
Single-duct air handler	14	14†	1 of 96
Fan-powered terminal units (parallel)	30	26	8 of 96
Fan-powered terminal units (series)	29	25	9 of 96
Dual-duct air handler	55	50	3 of 96
Fan coil unit	47	43	5 of 96
Rooftop unit (simulated)	24	20	0 of 28
Rooftop unit (measured)	1 case	—	3 of 49

A. Detection and false alarms

Across the five main systems the framework detected 158 of 175 seeded faults, with 26 false-alarm days in 480 held-out fault-free days. The median time from the start of a fault to its first alarm was one day. On the real rooftop unit the false-alarm rate was 3 of 49 held-out days. A later audit mapped every recorded column the first configurations had left unread; sixteen of the 158 detections came from columns that audit added. The 17 faults it missed are all coil fouling, where the recorded sensor values stay inside every healthy band. On the one dataset that records the refrigerant side, the framework caught every fouling severity.

B. Where the alarm points

On the two fan-powered systems, where each air handler serves four terminal units, 50 detections made before the column audit had their seeded fault in one terminal unit. In 37 of these the alarm named the correct terminal unit, in 13 it named the air handler but no terminal unit, and in none did it name a wrong one. Over all 158 detections the rule carrying the alarm pointed at the right class of component (coil, fan, damper or terminal unit) in 83 percent of cases and at the exact seeded component in 69 percent.

C. Against two baselines

Against principal component analysis, a standard statistical method, the framework detected 64 of 73 shared

TABLE II
STRATA against the Guideline 36 fault rules on the three systems on which both were run, under the same gate. "Detected" is seeded faults detected; "False" is false-alarm days among the held-out days. The single-duct count is the audited 13 of Table I.

	STRATA		Guideline 36 rules	
System	Detected	False	Detected	False
Single-duct air handler	13 of 14	1 of 96	12 of 14	1 of 96
Fan-powered terminal units (parallel)	26 of 30	8 of 96	8 of 30	0 of 96
Fan-powered terminal units (series)	25 of 29	9 of 96	9 of 29	0 of 96

scenarios against the baseline's 61, a difference too small to be statistically significant; the baseline cannot name the terminal unit. Against the Guideline 36 fault rules that commercial systems run, implemented by an independent open-source library and scored under the same gate, the framework detected one more fault than the rules on the single-duct air handler at the same false-alarm count, and three times as many on the terminal units, because it reads the per-zone sensors the rules do not (Table II). We had predicted that STRATA would detect at least as many faults as the rules at a lower false-alarm rate on every system. It did not: on the air handler the false-alarm counts were equal, and on the terminal units the extra coverage cost eight and nine false-alarm days against the rules' none.

D. What process mining contributes

We began with the hypothesis that the discovered process models would detect the faults themselves. They did not. Over 199 scenarios on six systems, no fault was caught by the conformance checks alone, and even on the scenario where the models did best, a one-line hand-written rule flagged more fault days than the model did. What process mining contributes is structure. The per-terminal-unit models give every alarm its location, and the models surfaced the regularities that became rules (Stage 3). On the two fan-powered systems the frequency check, which counts events in the log, is one of only two checks (with the oscillation check) that see the unstable-controller faults. Process mining helps to get detection but is not detecting on its own. That finding came from pre-registered tests, with the refuting result written down before each was run.

E. Robustness to noisier data

We applied four perturbations to the fault-free year and every fault file: sensor noise; change-of-value logging, which records a value only when it changes; gaps in the record; and shifted schedules. We then re-fitted and re-scored all five systems against a limit, set beforehand, of 10 false-alarm days in 96. Every system stayed under it, with the worst count at 7 of 96, and each condition cost at most one detection per system.

V. Deployment on a Real Building: The Next Step

None of this section has been carried out yet. Thompson Rivers University (TRU) is the planned site. A 2022 energy study of the Campus Activity Centre lists an air-handling unit serving variable-air-volume terminal units, fan coil units and five rooftop units [58], all equipment types on which the framework has been evaluated. Deployment has four steps. First, collect a fault-free period of trend data at the system's own sampling rate, a year if possible. Second, write the configuration file: every recorded point mapped to a role or excluded with a reason, and the control sequence as rules. Third, confirm silence: run the signature rules over the whole fault-free period and fix or remove any that fire, then calibrate the other checks and measure every check's false-alarm rate on this building's held-out days. Fourth, operate: run Stages 2 to 6 on every new day, with the gate's window set in the trial.

Three things decide whether it works. The fault-free period must really be fault-free, because one faulty day in it can move a band edge. The configuration must read every recorded point; the column audit of Section IV showed what unread columns cost. And the silence step must run at the building's own sampling rate, because thresholds set on one-minute data do not transfer to fifteen-minute data.

VI. Limitations

Every detection result above comes from simulated data, and simulated years are cleaner than real ones. Two scoring choices were made after we saw their effect: the rate check was demoted to support-only, and the two conformance checks were re-calibrated on days separate from those used to measure their false-alarm rates. Slowly developing faults are caught late, up to 114 days for two fouling cases. The benchmarks seed equipment and sensor faults, not the schedule and override faults that are most common in real buildings. Finally, the alarm names the symptom and the location, not the cause.

VII. Conclusion

STRATA turns the trend data a building already records into alarms an operator can act on. On the five main benchmark systems it detected 158 of 175 seeded faults and, in the 50 cases where this could be checked, never named a wrong terminal unit. The next step is a live trial on a real building, planned for the TRU campus.

Acknowledgment

The author thanks Dr. Anthony Aighobahi and Dr. Mridula Sharma for their supervision, and the Undergraduate Research Experience Award Program (U-REAP) for its support.

References

[1] United Nations Environment Programme and Global Alliance for Buildings and Construction, "Global status report for buildings and construction 2024/2025: Not just another brick in the wall," United Nations Environment Programme, Nairobi, Tech. Rep., 2025. [Online]. Available: https://globalabc.org/resources/publications/global-status-report-buildings-and-construction-20242025-not-just-another

[2] J. Bi, H. Wang, E. Yan, C. Wang, K. Yan, L. Jiang, and B. Yang, "AI in HVAC fault detection and diagnosis: A systematic review," Energy Reviews, vol. 3, no. 2, p. 100071, 2024.

[3] S. Katipamula and M. R. Brambley, "Methods for fault detection, diagnostics, and prognostics for building systems—a review, part I," HVAC&R Research, vol. 11, no. 1, pp. 3–25, 2005.

[4] E. Crowe, Y. Chen, H. Reeve, D. Yuill, A. Ebrahimifakhar, Y. Chen, L. Troup, A. Smith, and J. Granderson, "Empirical analysis of the prevalence of HVAC faults in commercial buildings," Science and Technology for the Built Environment, vol. 29, no. 10, pp. 1027–1038, 2023.

[5] W. M. P. van der Aalst, Process Mining: Data Science in Action, 2nd ed. Springer, 2016.

[6] J. Schein, S. T. Bushby, N. S. Castro, and J. M. House, "A rule-based fault detection method for air handling units," Energy and Buildings, vol. 38, no. 12, pp. 1485–1492, 2006.

[7] ASHRAE, "ASHRAE Guideline 36-2021: High-performance sequences of operation for HVAC systems," American Society of Heating, Refrigerating and Air-Conditioning Engineers, Peachtree Corners, GA, 2021. [Online]. Available: https://www.ashrae.org/technical-resources/ashrae-standards-and-guidelines

[8] ——, "ASHRAE Guideline 36-2024: High-performance sequences of operation for HVAC systems," American Society of Heating, Refrigerating and Air-Conditioning Engineers, Peachtree Corners, GA, 2024.

[9] J. Granderson, R. Singla, E. Mayhorn, P. Ehrlich, D. Vrabie, and S. Frank, "Characterization and survey of automated fault detection and diagnostic tools," Lawrence Berkeley National Laboratory, Berkeley, CA, Tech. Rep. LBNL-2001075, Nov. 2017. [Online]. Available: https://eta-publications.lbl.gov/sites/default/files/lbnl-2001075.pdf

[10] K. Heimar Andersen, S. Pommerencke Melgaard, H. Johra, A. Marszal-Pomianowska, R. Lund Jensen, and P. Kvols Heiselberg, "Barriers and drivers for implementation of automatic fault detection and diagnosis in buildings and HVAC systems: An outlook from industry experts," Energy and Buildings, vol. 303, p. 113801, 2024.

[11] G. Lin, H. Kramer, and J. Granderson, "Building fault detection and diagnostics: Achieved savings, and methods to evaluate algorithm performance," Building and Environment, vol. 168, p. 106505, 2020.

[12] X. Zhou, C. Battisti, R. Singla, G. LaPalme, G. Paliaga, and B. Eubanks, "Field demonstration of ASHRAE Guideline 36 high-performance sequences of operation for HVAC systems," Slipstream and TRC, for the Minnesota Department of Commerce, Conservation Applied Research and Development program, Tech. Rep., 2025. [Online]. Available: https://mn.gov/commerce-stat/energy/cards-2/187281/187281_slipstream_ashrae_guideline36_final_report_ada.pdf

[13] C. Battisti, Y. Li, A. Faramarzi, R. Singla, M. Mutmansky, M. Greene, and T. Jayarathne, "2028 California Energy Code CASE report: Nonresidential HVAC fault detection and diagnostics (draft)," California Statewide Codes and Standards Enhancement Team, Tech. Rep., 2026. [Online]. Available: https://title24stakeholders.com/wp-content/uploads/2026/03/2028_CASE-Report_Nonresidential-HVAC-Fault-Detection-and-Diagnostics.pdf

[14] A. R. Diallo, L. Homri, T. Boeuf, J.-Y. Dantan, and F. Bonnet, "Quantifying and mitigating alarm fatigue caused by fault detection systems," Reliability Engineering & System Safety, vol. 267, p. 111890, 2026.

[15] G. Lin, J. House, Y. Chen, J. Granderson, and H. Zhang, "Active multi-mode data analysis to improve fault diagnosis in AHUs," Energy and Buildings, vol. 337, p. 115621, 2025.

[16] A. Parsaei, A. A. Markus, B. Gunay, W. O'Brien, R. Moromisato, and J. Bursill, "A proactive fault detection and diagnostics method for zone-level auto-commissioning in VAV AHUs," Energy and Buildings, vol. 344, p. 116027, 2025.

[17] G. Lin, M. Pritoni, Y. Chen, and J. Granderson, "Development and implementation of fault-correction algorithms in fault detection and diagnostics tools," Energies, vol. 13, no. 10, p. 2598, 2020.

[18] M. Pritoni, G. Lin, Y. Chen, R. Vitti, C. Weyandt, and J. Granderson, "From fault-detection to automated fault correction: A field study," Building and Environment, vol. 214, p. 108900, 2022.

[19] Z. Wang, Y. Qin, Y. Kong, L. Wang, Q. Leng, and C. Zhang, "Advanced fault detection, diagnosis and prognosis in HVAC systems: Lifecycle insight, key challenges, and promising approaches," Renewable and Sustainable Energy Reviews, vol. 219, p. 115867, 2025.

[20] W. Tun, J. K.-W. Wong, and S.-H. Ling, "Hybrid random forest and support vector machine modeling for HVAC fault detection and diagnosis," Sensors, vol. 21, no. 24, p. 8163, 2021.

[21] Prince, B. Yoon, and P. Kumar, "Fault detection and diagnosis in air-handling unit (AHU) using improved hybrid 1D convolutional neural network," Systems, vol. 13, no. 5, p. 330, 2025.

[22] Q. Deng, Z. Chen, W. Zhu, Z. Li, Y. Yuan, and Y. Wang, "Adaptive fusion graph convolutional network based interpretable fault diagnosis method for HVAC systems enhanced by unlabeled data," Energy and Buildings, vol. 324, p. 114901, 2024.

[23] G. Li, Z. Yao, L. Chen, T. Li, and C. Xu, "An interpretable graph convolutional neural network based fault diagnosis method for building energy systems," Building Simulation, vol. 17, no. 7, pp. 1113–1136, 2024.

[24] H. Wang, W. Yan, Y. Zhang, J. Yu, W. Gao, B. Zhou, and Y. Liu, "xMHAC-PDCNN: An explainable deep network with spatial feature embedding for fault detection and diagnosis of chillers under variable operating conditions," Building and Environment, vol. 284, p. 113504, 2025.

[25] K.-S. Park, J.-H. Kim, and C. S. Park, "Explainable fault detection and diagnosis in HVAC systems using a Koopman operator with normal data," Energy and Buildings, vol. 350, p. 116580, 2026.

[26] M. Paolini, M. S. Piscitelli, and A. Capozzoli, "A label-free hybrid fault detection and diagnosis approach for HVAC systems using Bayesian networks," Energy and Buildings, vol. 351, p. 116658, 2026.

[27] M. A. F. Abdollah, R. Scoccia, and M. Aprile, "Transformer encoder based self-supervised learning for HVAC fault detection with unlabeled data," Building and Environment, vol. 258, p. 111568, 2024.

[28] K. El Mokhtari and J. J. McArthur, "Autoencoder-based fault detection using building automation system data," Advanced Engineering Informatics, vol. 62, p. 102810, 2024.

[29] X. Zhu, K. Chen, B. Anduv, X. Jin, and Z. Du, "Transfer learning based methodology for migration and application of fault detection and diagnosis between building chillers for improving energy efficiency," Building and Environment, vol. 200, p. 107957, 2021.

[30] N. Ghalamsiah, J. Wen, K. S. Candan, T. Wu, Z. O'Neill, and A. Aghaei, "Unsupervised domain adaptation for HVAC fault diagnosis using contrastive adaptation network," Energy and Buildings, vol. 337, p. 115659, 2025.

[31] S. Jiang, "Leakage-safe cross-domain AHU fault detection for building operation: adaptation reliability under limited target information," Energy and Buildings, vol. 372, p. 118275, 2026.

[32] S. P. Melgaard, R. L. Jensen, P. Ferreira, K. H. Andersen, N. Dionísio, V. V. Cogo, and P. Heiselberg, "Fault detection in AHU: A walkthrough for implementation in a Danish educational building," Energy and Buildings, vol. 342, p. 115844, 2025.

[33] A. Mukhtar, M. Hadwiger, F. Wotawa, and G. Schweiger, "Reproducibility of machine learning-based fault detection and

diagnosis for HVAC systems in buildings: An empirical study," Energy and AI, vol. 22, p. 100658, 2025.
[34] S. Mazzetto, "Hybrid predictive maintenance for building systems: Integrating rule-based and machine learning models for fault detection using a high-resolution Danish dataset," Buildings, vol. 15, no. 4, p. 630, 2025.
[35] H. R. Y. Zabadi, F. A. Zemedkun, H. Hosamo, and D. Kraniotis, "Integration of digital twins and machine learning for predictive maintenance using APAR method rules in non-residential buildings," Frontiers in Built Environment, vol. 11, p. 1734945, 2026.
[36] S. Lin and C. Hua, "Physics-informed large language models for HVAC anomaly detection with autonomous rule generation," arXiv preprint arXiv:2510.17146, 2025. [Online]. Available: https://arxiv.org/abs/2510.17146
[37] A. Neubauer, T. Hong, H. Li, M. Yu, A. Darbandi, Y. Fürst, and M. Kriegel, "Large language models for HVAC operations in building energy systems: A critical review of methods, applications, and deployment readiness," arXiv preprint arXiv:2609.05314, 2026. [Online]. Available: https://arxiv.org/abs/2609.05314
[38] S. J. J. Leemans, D. Fahland, and W. M. P. van der Aalst, "Discovering block-structured process models from event logs — a constructive approach," in Application and Theory of Petri Nets and Concurrency (PETRI NETS 2013), ser. Lecture Notes in Computer Science, vol. 7927. Springer, 2013, pp. 311–329.
[39] J. Carmona, B. van Dongen, A. Solti, and M. Weidlich, Conformance Checking: Relating Processes and Models. Springer, 2018.
[40] F. Bezerra and J. Wainer, "Algorithms for anomaly detection of traces in logs of process aware information systems," Information Systems, vol. 38, no. 1, pp. 33–44, 2013.
[41] D. Janssen, F. Mannhardt, A. Koschmider, and S. J. van Zelst, "Process model discovery from sensor event data," in Process Mining Workshops (ICPM 2020), ser. Lecture Notes in Business Information Processing, vol. 406. Springer, 2021, pp. 69–81.
[42] E. Brzychczy, M. Aleknonytė-Resch, D. Janssen, and A. Koschmider, "Process mining on sensor data: a review of related works," Knowledge and Information Systems, vol. 67, pp. 4915–4948, 2025.
[43] T. Tratnjek, M. Kocbek Bule, and G. Polančič, "From signals to events: A maturity-based perspective on IoT-enabled process mining," Sensors, vol. 26, no. 19, p. 6293, 2026.
[44] C. Imenkamp, H. Kirchmann, S. Plumhoff, M. Weidlich, and A. Koschmider, "Glass-box activity discovery from raw sensor data," BP-Meet-IoT 2026 workshop, research track, 2026. [Online]. Available: https://bp-meet-iot.github.io/papers/bp-meet-iot2026-activity-discovery.pdf
[45] Y. Shi, N. Zhang, X. Song, H. Li, and Q. Zhu, "Novel approach for industrial process anomaly detection based on process mining," Journal of Process Control, vol. 136, p. 103165, 2024.
[46] A. Norouzifar, A. Neshastegaran, and I. Izadi, "Modeling alarm sequences using process mining techniques," in 2026 34th International Conference on Electrical Engineering (ICEE), 2026, pp. 1428–1433.
[47] E. J. Husom, A. Goknil, F. Mannhardt, S. Tverdal, S. Sen, and P. Nguyen, "Unsupervised learning and process analysis for sensor data validation in the IIoT," ACM Transactions on Internet Technology, vol. 26, no. 2, pp. 1–34, 2026.
[48] F. Vitale, S. Guarino, F. Flammini, L. Faramondi, N. Mazzocca, and R. Setola, "Process mining for digital twin development of industrial cyber-physical systems," IEEE Transactions on Industrial Informatics, vol. 21, pp. 866–875, 2025.
[49] F. Vitale, N. Dall'Ora, S. Gaiardelli, E. Fraccaroli, N. Mazzocca, and F. Fummi, "Process mining-driven modeling and simulation to enhance fault diagnosis in cyber-physical systems," Journal of Manufacturing Systems, vol. 84, pp. 189–206, 2026.
[50] F. Vitale, M. Pegoraro, W. M. P. van der Aalst, and N. Mazzocca, "Control-flow anomaly detection by process mining-based feature extraction and dimensionality reduction," Knowledge-Based Systems, vol. 310, p. 112970, 2025.
[51] A. Moradbeikie, A. P. A. d. C. Barbon, I. M. Grigore, D. F. Barbin, and S. Barbon Junior, "Process mining of sensor data for predictive process monitoring: A HACCP-guided pasteurization study case," Systems, vol. 13, no. 11, p. 935, 2025.
[52] W. Deabes, K. E. Bouazza, and W. Algthami, "Smart fuzzy Petri net-based temperature control framework for reducing building energy consumption," Sensors, vol. 23, no. 13, p. 5985, 2023.
[53] M. Wetter, P. Ehrlich, A. Gautier, M. Grahovac, P. Haves, J. Hu, A. Prakash, D. Robin, and K. Zhang, "OpenBuildingControl: Digitizing the control delivery from building energy modeling to specification, implementation and formal verification," Energy, vol. 238, p. 121501, 2022.
[54] K. Akhramovich, E. Serral, and C. Cetina, "A systematic literature review on the application of process mining to Industry 4.0," Knowledge and Information Systems, vol. 66, no. 5, pp. 2699–2746, 2024.
[55] F. De Luzi, F. Leotta, A. Marrella, and M. Mecella, "On the interplay between business process management and Internet-of-Things: A systematic literature review," Business & Information Systems Engineering, vol. 67, no. 2, pp. 265–288, 2025.
[56] Z. Khais Shahid, S. Saguna, and C. Åhlund, "Unsupervised analysis of EcoStruxure Building Operation (EBO) alarm logs in school buildings," in Proceedings of the 13th ACM International Conference on Systems for Energy-Efficient Buildings, Cities, and Transportation (BuildSys '26), 2026, pp. 282–292.
[57] J. Granderson, G. Lin, Y. Chen, A. Casillas, J. Wen, Z. Chen, P. Im, S. Huang, and J. Ling, "A labeled dataset for building HVAC systems operating in faulted and fault-free states," Scientific Data, vol. 10, p. 342, 2023.
[58] SES Consulting Inc., "Thompson Rivers University ASHRAE Level 1 energy study: Campus Activity Centre," Prepared for Thompson Rivers University, Vancouver, BC, Tech. Rep., May 2022. [Online]. Available: https://www.tru.ca/___shared/assets/level-1-energy-study-2022-cac64231.pdf